

Foundations for statistical inference - Confidence intervals

Sampling from Ames, Iowa

If you have access to data on an entire population, say the size of every house in Ames, Iowa, it's straight forward to answer questions like, "How big is the typical house in Ames?" and "How much variation is there in sizes of houses?". If you have access to only a sample of the population, as is often the case, the task becomes more complicated. What is your best guess for the typical size if you only know the sizes of several dozen houses? This sort of situation requires that you use your sample to make inference on what your population looks like.

The data

In the previous lab, "Sampling Distributions", we looked at the population data of houses from Ames, Iowa. Let's start by loading that data set.

```
load("more/ames.RData")
```

In this lab we'll start with a simple random sample of size 60 from the population. Specifically, this is a simple random sample of size 60. Note that the data set has information on many housing variables, but for the first portion of the lab we'll focus on the size of the house, represented by the variable `Gr.Liv.Area`.


```
population <- ames$Gr.Liv.Area  
samp <- sample(population, 60)
```

1. Describe the distribution of your sample. What would you say is the "typical" size within your sample? Also state precisely what you interpreted "typical" to mean.

```
summary(samp)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   
##      720    1157    1655    1546    1852    2630
```

```
hist(samp,breaks=15)
```



The interquartile range here is 1139 - 1828. With a median of 1580 and a mean of 156.

The data is very roughly normally distributed according to the histogram, so I'd say the "typical" size would be one which is in the IQR. This would include 50% of our sample.

2. Would you expect another student's distribution to be identical to yours? Would you expect it to be similar? Why or why not?

I would expect another student's distribution to be similar to mine, but not identical. The nature of random sampling leads to different results, but they should be close to the population distribution

Confidence intervals

One of the most common ways to describe the typical or central value of a distribution is to use the mean. In this case we can calculate the mean of the sample using,

```
sample_mean <- mean(samp)
```

Return for a moment to the question that first motivated this lab: based on this sample, what can we infer about the population? Based only on this single sample, the best estimate of the average living area of houses sold in Ames would be the sample mean, usually denoted as \bar{x} (here we're calling it `sample_mean`). That serves as a good *point estimate* but it would be useful to also communicate how uncertain we are of that estimate. This can be captured by using a *confidence interval*.

We can calculate a 95% confidence interval for a sample mean by adding and subtracting 1.96 standard errors to the point estimate (See Section 4.2.3 if you are unfamiliar with this formula).

```
se <- sd(samp) / sqrt(60)
lower <- sample_mean - 1.96 * se
upper <- sample_mean + 1.96 * se
c(lower, upper)
```

```
## [1] 1423.734 1668.932
```

This is an important inference that we’ve just made: even though we don’t know what the full population looks like, we’re 95% confident that the true average size of houses in Ames lies between the values *lower* and *upper*. There are a few conditions that must be met for this interval to be valid.

3. For the confidence interval to be valid, the sample mean must be normally distributed and have standard error s/\sqrt{n} . What conditions must be met for this to be true?

According to the central limit theorem, if a sample consists of at least 30 independent observations, and the data are not strongly skewed, then the distribution of the sample mean is well approximated by the normal model. This is true for any population distribution.

Confidence levels

4. What does “95% confidence” mean? If you’re not sure, see Section 4.2.2.

The 95% confidence interval can be interpreted in terms of repeat samples. If you take multiple samples, 95% of the samples can be expected to contain the actual mean.

In this case we have the luxury of knowing the true population mean since we have data on the entire population. This value can be calculated using the following command:

```
mean(population)
```

```
## [1] 1499.69
```

5. Does your confidence interval capture the true average size of houses in Ames? If you are working on this lab in a classroom, does your neighbor’s interval capture this value?

My confidence interval (at least initially calculated) did indeed capture the population mean.

6. Each student in your class should have gotten a slightly different confidence interval. What proportion of those intervals would you expect to capture the true population mean? Why? If you are working in this lab in a classroom, collect data on the intervals created by other students in the class and calculate the proportion of intervals that capture the true population mean.

95% of the intervals calculated in the class should capture the mean, because we’re using a z-score of ± 1.96 to calculate them.

Using R, we’re going to recreate many samples to learn more about how sample means and confidence intervals vary from one sample to another. *Loops* come in handy here (If you are unfamiliar with loops, review the [Sampling Distribution Lab](#)).

Here is the rough outline:

- Obtain a random sample.
- Calculate and store the sample's mean and standard deviation.
- Repeat steps (1) and (2) 50 times.
- Use these stored statistics to calculate many confidence intervals.

But before we do all of this, we need to first create empty vectors where we can save the means and standard deviations that will be calculated from each sample. And while we're at it, let's also store the desired sample size as `n`.

```
samp_mean <- rep(NA, 50)
samp_sd <- rep(NA, 50)
n <- 60
```

Now we're ready for the loop where we calculate the means and standard deviations of 50 random samples.

```
for(i in 1:50){
  samp <- sample(population, n) # obtain a sample of size n = 60 from the population
  samp_mean[i] <- mean(samp)    # save sample mean in ith element of samp_mean
  samp_sd[i] <- sd(samp)        # save sample sd in ith element of samp_sd
}
```

Lastly, we construct the confidence intervals.

```
lower_vector <- samp_mean - 1.96 * samp_sd / sqrt(n)
upper_vector <- samp_mean + 1.96 * samp_sd / sqrt(n)
```

Lower bounds of these 50 confidence intervals are stored in `lower_vector`, and the upper bounds are in `upper_vector`. Let's view the first interval.


```
c(lower_vector[1], upper_vector[1])
```

```
## [1] 1254.188 1504.845
```

On your own

- Using the following function (which was downloaded with the data set), plot all intervals. What proportion of your confidence intervals include the true population mean? Is this proportion exactly equal to the confidence level? If not, explain why.

```
plot_ci(lower_vector, upper_vector, mean(population))
```



- Pick a confidence level of your choosing, provided it is not 95%. What is the appropriate critical value?
Let's choose an 85% confidence interval.

I'll use the `qnorm` function to calculate this, but I can't just enter `qnorm(0.85)`. 1.96 is a critical range, meaning if we have a normal distribution $N(\mu = 0, \sigma = 1)$,

This means that there's 2.5% under the tails beyond 1.96. For an 85% confidence interval, there's 15% under both tails, and 7.5% under the right. So, I'll use `qnorm` for `1-0.075`

```
criticalvalue85 <- qnorm(1-0.075)
```


```
criticalvalue85
```

```
## [1] 1.439531
```

- Calculate 50 confidence intervals at the confidence level you chose in the previous question. You do not need to obtain new samples, simply calculate new intervals based on the sample means and standard deviations you have already collected. Using the `plot_ci` function, plot all intervals and calculate the proportion of intervals that include the true population mean. How does this percentage compare to the confidence level selected for the intervals?

```
lower_vector <- samp_mean - criticalvalue85 * samp_sd / sqrt(n)
upper_vector <- samp_mean + criticalvalue85 * samp_sd / sqrt(n)
```

```
plot_ci(lower_vector, upper_vector, mean(population))
```



This makes sense. At each sample's mean, we're calculating an interval that we're 85% sure the mean resides in. It will be a shorter interval, but this means that less samples will include the population mean.

This is a product of OpenIntro that is released under a [Creative Commons Attribution-ShareAlike 3.0 Unported](#). This lab was written for OpenIntro by Andrew Bray and Mine Çetinkaya-Rundel.